

Exploration

```
import pandas as pd
import matplotlib.pyplot as plt

trainingSet = pd.read_csv("./data/train.csv")
testingSet = pd.read_csv("./data/test.csv")

print("train.csv shape is ", trainingSet.shape)
print("test.csv shape is ", testingSet.shape)

print()

print(trainingSet.head())
print()
print(testingSet.head())

print()

print(trainingSet.describe())

trainingSet['Score'].value_counts().plot(kind='bar', legend=True,
alpha=.5)
plt.title("Count of Scores")
plt.show()

trainingSet['ProductId'].value_counts().nlargest(25).plot(kind='bar',
legend=True, alpha=.5)
plt.title("Top 25 most rated Products")
plt.show()

trainingSet['ProductId'].value_counts().nsmallest(25).plot(kind='bar',
legend=True, alpha=.5)
plt.title("Top 25 least rated Products")
plt.show()

trainingSet['UserId'].value_counts().nlargest(25).plot(kind='bar',
legend=True, alpha=.5)
plt.title("Top 25 Reviewers")
plt.show()

trainingSet['UserId'].value_counts().nsmallest(25).plot(kind='bar',
legend=True, alpha=.5)
plt.title("Lowest 25 Reviewers")
plt.show()

trainingSet[['Score',
'HelpfulnessNumerator']].groupby('Score').mean().plot(kind='bar',
legend=True, alpha=.5)
plt.title("Mean Helpfulness Numerator per Score")
```

```

plt.show()

trainingSet[['Score',
'ProductId']].groupby('ProductId').mean().nlargest(25,
'Score').plot(kind='bar', legend=True, alpha=.5)
plt.title("Top 25 best rated Products")
plt.show()

trainingSet[['Score',
'ProductId']].groupby('ProductId').mean().nsmallest(25,
'Score').plot(kind='bar', legend=True, alpha=.5)
plt.title("Top 25 worst rated Products")
plt.show()

trainingSet[['Score', 'UserId']].groupby('UserId').mean().nlargest(25,
'Score').plot(kind='bar', legend=True, alpha=.5)
plt.title("Top 25 kindest Reviewers")
plt.show()

trainingSet[['Score',
'UserId']].groupby('UserId').mean().nsmallest(25,
'Score').plot(kind='bar', legend=True, alpha=.5)
plt.title("Top 25 harshest Reviewers")
plt.show()

trainingSet[trainingSet['ProductId'].isin(trainingSet['ProductId'].value_counts().nlargest(25).index.tolist())][['Score',
'ProductId']].groupby('ProductId').mean().plot(kind='bar',
legend=True, alpha=.5)
plt.title("Mean of top 25 most rated Products")
plt.show()

train.csv shape is (139753, 9)
test.csv shape is (17470, 2)

```

	Id	ProductId	UserId	HelpfulnessNumerator \
0	195370	1890228583	A3VLX5Z090RQ0V	1
1	1632470	B00BEIYSL4	AUDXDMFM49NGY	0
2	9771	0767809335	A3LFIA97BUU5IE	3
3	218855	6300215792	A1QZM75342ZQVQ	1
4	936225	B000B5X0ZW	ANM2SCEUL3WL1	1

	HelpfulnessDenominator	Time \
0	2	1030838400
1	1	1405036800
2	36	983750400
3	1	1394841600
4	1	1163721600

Summary \

```

0           An Unexplained Anime Review
1                               not great.
2           Technical problem with this DVD
3           Heeeeyyyy LAAAAADEEEE!!!!
4 Herzog the Great Traveler of both natural and ...


```

	Text	Score
0	I was very anxious to see the Uncut version of...	2.0
1	Movie was okay...not great.	3.0
2	Like the Dinosaur Collector's Edition DVD, thi...	1.0
3	Come on, now..... this has to be, by far, the...	5.0
4	I've always been a great admirer of Herzog's o...	4.0

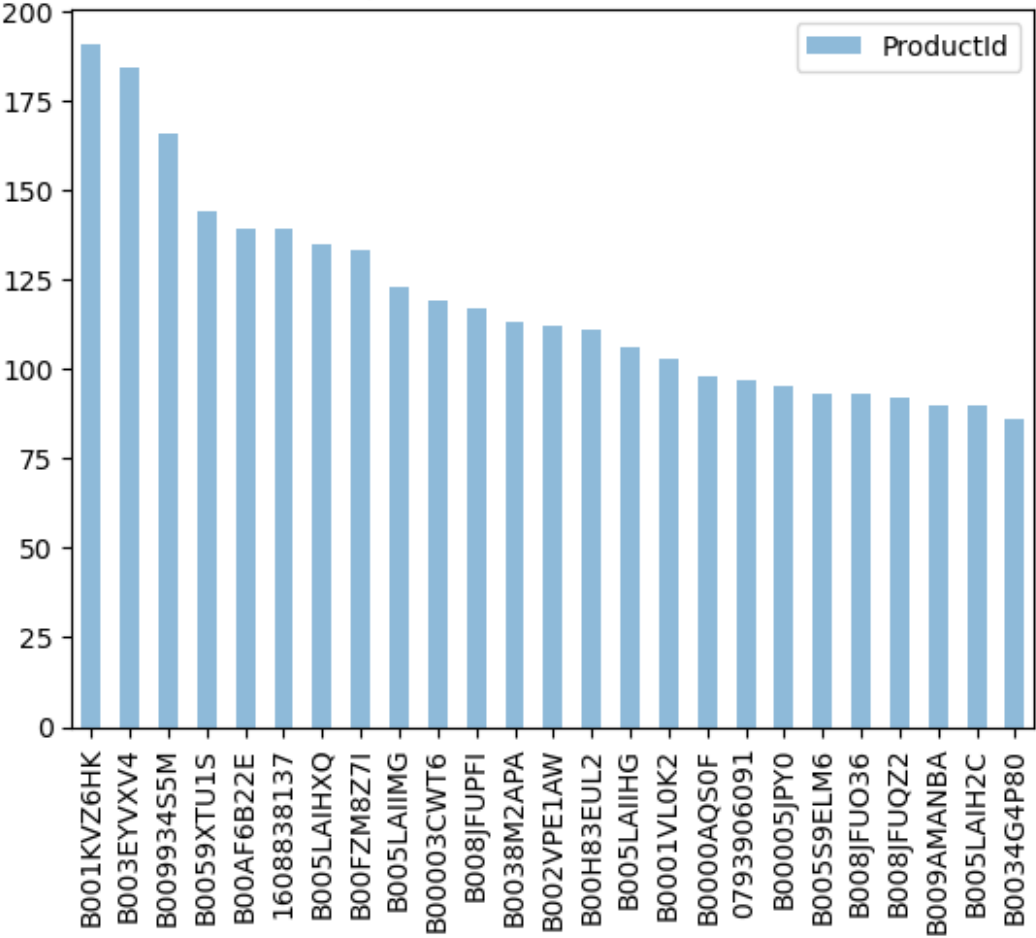
	Id	Score
0	786781	NaN
1	17153	NaN
2	1557328	NaN
3	1242666	NaN
4	1359242	NaN

	Id	HelpfulnessNumerator	HelpfulnessDenominator \
count	1.397530e+05	139753.000000	139753.000000
mean	8.497881e+05	3.601096	5.313246
std	4.896942e+05	20.101195	22.300962
min	8.000000e+00	0.000000	0.000000
25%	4.258660e+05	0.000000	0.000000
50%	8.510200e+05	1.000000	1.000000
75%	1.273392e+06	3.000000	5.000000
max	1.697519e+06	4646.000000	4682.000000


	Time	Score
count	1.397530e+05	122283.000000
mean	1.262516e+09	4.115552
std	1.287262e+08	1.191661
min	8.948448e+08	1.000000
25%	1.164758e+09	4.000000
50%	1.307318e+09	5.000000
75%	1.373155e+09	5.000000
max	1.406074e+09	5.000000




Top 25 most rated Products




Top 25 least rated Products




Top 25 Reviewers




Lowest 25 Reviewers







Top 25 best rated Products




Top 25 worst rated Products

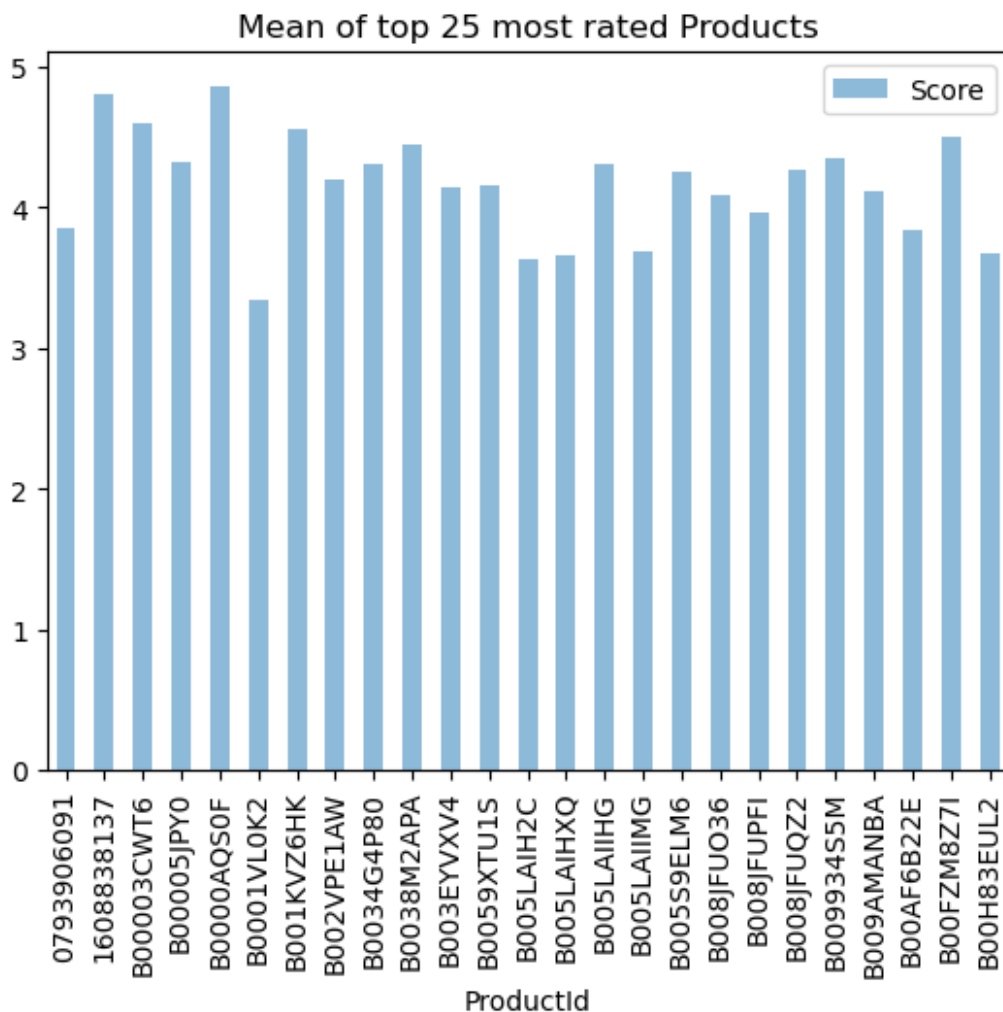


Top 25 kindest Reviewers



Top 25 harshest Reviewers





Feature Extraction

```
import pandas as pd

def process(df):
    # This is where you can do all your processing

    df['Helpfulness'] = df['HelpfulnessNumerator'] /
df['HelpfulnessDenominator']
    df['Helpfulness'] = df['Helpfulness'].fillna(0)

    df['ReviewLength'] = df.apply(lambda row :
len(row['Text'].split()) if type(row['Text']) == str else 0, axis = 1)

    return df

# Load the dataset
trainingSet = pd.read_csv("./data/train.csv")
```

```

# Process the DataFrame
train_processed = process(trainingSet)

# Load test set
submissionSet = pd.read_csv("./data/test.csv")

# Merge on Id so that the test set can have feature columns as well
testX= pd.merge(train_processed, submissionSet, left_on='Id',
right_on='Id')
testX = testX.drop(columns=['Score_x'])
testX = testX.rename(columns={'Score_y': 'Score'})

# The training set is where the score is not null
trainX = train_processed[train_processed['Score'].notnull()]

# Save the datasets with the new features for easy access later
testX.to_csv("./data/X_test.csv", index=False)
trainX.to_csv("./data/X_train.csv", index=False)

```

Creating your model

```

import pickle
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, confusion_matrix,
mean_squared_error

# Load training set with new features into DataFrame
X_train = pd.read_csv("./data/X_train.csv")

# Split training set into training and testing set
X_train, X_test, Y_train, Y_test = train_test_split(
    X_train.drop(['Score'], axis=1),
    X_train['Score'],
    test_size=1/4.0,
    random_state=3
)

# This is where you can do more feature selection
X_train_processed = X_train.drop(columns=['Id', 'ProductId', 'UserId',
'Text', 'Summary'])
X_test_processed = X_test.drop(columns=['Id', 'ProductId', 'UserId',
'Text', 'Summary'])

# Learn the model
model = KNeighborsClassifier(n_neighbors=20).fit(X_train_processed,
Y_train)

```

```


# pickle model - saves it so you can load it later
with open('knn_20_model.obj', 'wb') as f:
    pickle.dump(model, f)
# to load pickled model:
# with open('filename', 'rb') as f:
#     model = pickle.load(f)

# Evaluate your model on the testing set
Y_test_predictions = model.predict(X_test_processed)
print("Accuracy on testing set = ", accuracy_score(Y_test,
Y_test_predictions))
print("RMSE on testing set = ", mean_squared_error(Y_test,
Y_test_predictions) ** 0.5)

# Plot a confusion matrix
cm = confusion_matrix(Y_test, Y_test_predictions, normalize='true')
sns.heatmap(cm, annot=True)
plt.title('Confusion matrix of the classifier')
plt.xlabel('Predicted')
plt.ylabel('True')
plt.show()

Accuracy on testing set = 0.5121520395145727
RMSE on testing set = 1.4600992087486073

```

Create the Kaggle submission

```
X_submission = pd.read_csv("./data/X_test.csv")
X_submission_processed = X_submission.drop(columns=['Id', 'ProductId',
'UserId', 'Text', 'Summary', 'Score'])

X_submission['Score'] = model.predict(X_submission_processed)
submission = X_submission[['Id', 'Score']]
submission.to_csv("./data/submission.csv", index=False)
```

Now you can upload the `submission.csv` to kaggle